

Figure 2-3. CPU memory map.

2.3 Registers

The 6502 has fewer registers than similar processors. There are three special purpose registers, the program counter, stack pointer and status register which each have a specific use. It also has three general purpose registers, the accumulator and the index registers, X and Y, which can be used to store data or control information temporarily.

2.3.1 Program Counter (PC)

The program counter is a 16-bit register which holds the address of the next instruction to be executed. As instructions are executed, the value of the program counter is updated, usually moving on to the next instruction in the sequence. The value can be affected by branch and jump instructions, procedure calls and interrupts.

2.3.2 Stack Pointer (SP)

The stack is located at memory locations \$0100-\$01FF. The stack pointer is an 8-bit register which serves as an offset from \$0100. The stack works top-down, so when a byte is pushed on to the stack, the stack pointer is decremented and when a byte is pulled from the stack, the stack pointer is incremented. There is no detection of stack overflow and the stack pointer will just wrap around from \$00 to \$FF.

2.3.3 Accumulator (A)

The accumulator is an 8-bit register which stores the results of arithmetic and logic operations. The accumulator can also be set to a value retrieved from memory.

2.3.4 Index Register X (X)

The X register is an 8-bit register typically used as a counter or an offset for certain addressing modes. The X register can be set to a value retrieved from memory and can be used to get or set the value of the stack pointer.

2.2.5 Index Register Y (Y)

The Y register is an 8-bit register which is used in the same way as the X register, as a counter or to store an offset. Unlike the X register, the Y register cannot affect the stack pointer.

2.3.6 Processor Status (P)

The status register contains a number of single bit flags which are set or cleared when instructions are executed.

- Carry Flag (C) - The carry flag is set if the last instruction resulted in an overflow from bit 7 or an underflow from bit 0. For example performing $255 + 1$ would result in an answer of 0 with the carry bit set. This allows the system to perform calculations on numbers longer than 8-bits by performing the calculation on the first byte, storing the carry and then using that carry when performing the calculation on the second byte. The carry flag can be set by the SEC (Set Carry Flag) instruction and cleared by the CLC (Clear Carry Flag) instruction.
- Zero Flag (Z) - The zero flag is set if the result of the last instruction was zero. So for example $128 - 127$ does not set the zero flag, whereas $128 - 128$ does.
- Interrupt Disable (I) - The interrupt disable flag can be used to prevent the system responding to IRQs. It is set by the SEI (Set Interrupt Disable) instruction and IRQs will then be ignored until execution of a CLI (Clear Interrupt Disable) instruction.
- Decimal Mode (D) - The decimal mode flag is used to switch the 6502 into BCD mode. However the 2A03 does not support BCD mode so although the flag can be set, its value will be ignored. This flag can be set SED (Set Decimal Flag) instruction and cleared by CLD (Clear Decimal Flag).
- Break Command (B) - The break command flag is used to indicate that a BRK (Break) instruction has been executed, causing an IRQ.
- Overflow Flag (V) - The overflow flag is set if an invalid two's complement result was obtained by the previous instruction. This means that a negative result has been obtained when a positive one was expected or vice versa. For example, adding two positive numbers should give a positive answer. However $64 + 64$ gives the result -128 due to the sign bit. Therefore the overflow flag would be set. The overflow flag is determined by taking the exclusive-or of the carry from between bits 6 and 7 and between bit 7 and the carry flag [29]. An explanation of two's complement can be found in Appendix A.

- Negative Flag (N) - Bit 7 of a byte represents the sign of that byte, with 0 being positive and 1 being negative. The negative flag (also known as the sign flag) is set if this sign bit is 1.

The flags are arranged in the status register in the order shown in figure 2-3. Bit 5 of the status register is unused.




Figure 2-4. Status register layout.

2.4 Interrupts

Interrupts prevent the standard sequential execution of code and cause the processor to attend to the interrupt. Interrupts are usually generated by hardware which requires attention, but can be triggered by software. The NES has three different types of interrupt, NMI, IRQ and reset. The addresses to jump to when an interrupt occurs are stored in a vector table in the program code at \$FFFA-\$FFFF. When an interrupt occurs the system performs the following actions [30]:

1. Recognize interrupt request has occurred.
2. Complete execution of the current instruction.
3. Push the program counter and status register on to the stack.
4. Set the interrupt disable flag to prevent further interrupts.
5. Load the address of the interrupt handling routine from the vector table into the program counter.
6. Execute the interrupt handling routine.
7. After executing a RTI (Return From Interrupt) instruction, pull the program counter and status register values from the stack.
8. Resume execution of the program.

IRQs, or maskable interrupts, are generated by certain memory mappers. They are ignored by the processor if the interrupt disable flag is set. IRQs can be triggered by the software by use of the BRK (Break) instruction. When an IRQ occurs the system jumps to the address located at \$FFFE and \$FFFF.

NMI (Non-Maskable Interrupt) is the type of interrupt generated by the PPU when V-Blank occurs at the end of each frame. NMIs are not affected by the interrupt disable bit in the status register, so execution is always interrupted when they occur [31]. However, triggering of a NMI can be prevented if bit 7 of PPU Control Register 1 (\$2000) is clear. When a NMI occurs the system jumps to the address located at \$FFFA and \$FFFB. The handling of NMIs is shown in figure 2-4.

Reset interrupts are triggered when the system first starts and when the user presses the reset button. When a reset occurs the system jumps to the address located at \$FFFC and \$FFFD.

The system gives the highest priority to reset requests, followed by NMI and finally IRQ [7]. The NES has an interrupt latency of 7 cycles, which means it takes 7 CPU cycles to begin executing the interrupt handler.




Figure 2-5. NMI (Non-Maskable Interrupt) handling.

2.5 Addressing Modes

The 6502 has several different addressing modes, providing different ways to access memory locations. There are also addressing modes which operate on the contents of registers, rather than memory. In total there are 13 different addressing modes on the 6502. Some instructions can use more than one different addressing mode. Details on the available addressing modes can be found in Appendix E.

2.6 Instructions

The 6502 has 56 different instructions although some come in multiple variations using different addressing modes, making a total of 151 valid opcodes (operation codes) out of a possible 256. A detailed explanation of the complete instruction set can be found in [2], [29] and [32]. Instructions are either one, two or three bytes long, depending on the addressing mode. The first byte is the opcode and the remaining bytes are the operands. Instructions fit into several functional groups [3]:

- Load / Store Operations - Load a register from memory or stores the contents of a register to memory.
- Register Transfer Operations - Copy contents of X or Y register to the accumulator or copy contents of accumulator to X or Y register.
- Stack Operations - Push or pull the stack or manipulate stack pointer using X register.
- Logical Operations - Perform logical operations on the accumulator and a value stored in memory.
- Arithmetic Operations - Perform arithmetic operations on registers and memory.
- Increments / Decrements - Increment or decrement the X or Y registers or a value stored in memory.
- Shifts - Shift the bits of either the accumulator or a memory location one bit to the left or right.

- Jumps / Calls - Break sequential execution sequence, resuming from a specified address.
- Branches - Break sequential execution sequence, resuming from a specified address, if a condition is met. The condition involves examining a specific bit in the status register.
- Status Register Operations - Set or clear a flag in the status register.
- System Functions - Perform rarely used functions.

3 - Picture Processing Unit

3.1 2C02 Overview

Ricoh also supplied the 2C02 to serve as PPU. The PPU's registers are mostly located in the I/O registers section of CPU memory at \$2000-\$2007 and \$4014 as described in Appendix B. In addition, there are some special registers used for screen scrolling.

3.2 PPU Memory Map

The PPU has its own memory, known as VRAM (Video RAM). Like the CPU, the PPU can also address 64 KB of memory although it only has 16 KB of physical RAM. The PPU's memory map is shown in figure 3-1. Again, the left hand map shows a simplified version which is elaborated on by the right hand map. Due to the difference between physical and logical address spaces, any address above \$3FFF is wrapped around, making the logical memory locations \$4000-\$FFFF effectively a mirror of locations \$0000-\$3FFF.

Reading from and writing to PPU memory can be done by using the I/O registers \$2006 and \$2007 in CPU memory. This is usually done during V-Blank at the end of a frame, as it affects addresses used while drawing the screen and can therefore corrupt what is displayed. However, this effect can be used to produce split screen effects.

Since PPU memory uses 16-bit addresses but I/O registers are only 8-bit, two writes to \$2006 are required to set the address required. Data can then be read from or written to \$2007. After each write to \$2007, the address is incremented by either 1 or 32 as dictated by bit 2 of \$2000. The first read from \$2007 is invalid and the data will actually be buffered and returned on the next read. This does not apply to colour palettes.

The PPU also has a separate 256 byte area of memory, SPR-RAM (Sprite RAM), to store the sprite attributes. The sprites themselves can be found in the pattern tables.




Figure 3-1. PPU memory map.

3.3 PPU Registers

Communication between the CPU and other devices takes place via memory mapped I/O registers. The registers used by the PPU are located in main memory at \$2000-\$2007 with an additional register used for Direct Memory Access at \$4014. Remember that locations \$2000-\$2007 are mirrored every 8 bytes in the region \$2008-\$3FFF. A summary of all I/O registers can be found in Appendix B.

The actions of the PPU can be controlled by the CPU by writing to \$2000 and \$2001, known as PPU Control Register 1 and PPU Control Register 2 respectively. Both registers should

only be written to. Bit 7 of \$2000 can be used to disable NMIs. Remember that this type of interrupt is generated whenever a V-Blank occurs and is unaffected by the interrupt disable flag of the status register. Clearing this bit will prevent an NMI from occurring on V-Blank. Since the NES supports both 8x8 and 8x16 sprites, setting bit 5 of \$2000 will switch to 8x16 sprites. The next address in PPU memory to read or write from will be incremented after each I/O occurs. The value to increment by is adjusted by setting the value of bit 2 of \$2000. If this is clear, the address is incremented by 1 (horizontal), otherwise the increment is 32 (vertical). Using \$2001, the background can be hidden by clearing bit 3 and, similarly, the sprites can be hidden by clearing bit 4.

The PPU Status Register is located at \$2002 and is read only. The register is used by the PPU to report its status to the CPU. The programs will frequently cause the CPU to read from this location in order to ascertain the PPU's status. Bit 7 is set by the PPU to indicate that V-Blank is occurring. Bit 6 and bit 7 relate to sprites and are described later. Bit 4 indicates whether the PPU is willing to accept writes to VRAM, if it clear then writes are ignored. When a read from \$2002 occurs, bit 7 is reset to 0 as are \$2005 and \$2006.

3.3.1 Direct Memory Access

When transferring a large amount of data between devices it is inefficient to transfer this through the processor. To transfer data from CPU memory to sprite memory, for example, takes the following steps:

1. Load required SPR-RAM address into CPU.
2. Write required SPR-RAM address to \$2003.
3. Load byte into CPU.
4. Write byte to \$2004.

When filling the contents of sprite memory, this technique would have to be repeated 256 times. Direct Memory Access (DMA) is a technique which allows more efficient copying of data from CPU memory to sprite memory. Using DMA, the whole of sprite memory can be filled by using a single instruction, a write to \$4014. The starting address in CPU memory is specified by the operand for the write multiplied by \$100. The 256 bytes starting at this address are copied directly into sprite memory without the further involvement of the CPU.

When the DMA is occurring, the memory bus is in use, preventing the CPU from accessing memory and, therefore, preventing it from accessing any more instructions. This is referred to as cycle stealing and the CPU has to wait until the DMA transfer is complete. On the NES, the DMA takes the equivalent of 512 cycles (about 4.5 scanlines worth) after which the CPU can resume. This is considerably less than would be required to copy manually through the CPU.

3.4 Colour Palette

The NES has a colour palette containing 52 colours although there is actually room for 64. However, not all of these can be displayed at a given time. The NES uses two palettes, each with 16 entries, the image palette (\$3F00-\$3F0F) and the sprite palette (\$3F10-\$3F1F). The image palette shows the colours currently available for background tiles. The sprite palette shows the colours currently available for sprites. These palettes do not store the actual colour values but rather the index of the colour in the system palette. Since only 64 unique values are needed, bits 6 and 7 can be ignored.

The palette entry at \$3F00 is the background colour and is used for transparency. Mirroring is used so that every four bytes in the palettes is a copy of \$3F00. Therefore \$3F04, \$3F08, \$3F0C, \$3F10, \$3F14, \$3F18 and \$3F1C are just copies of \$3F00 and the total number of

colours in each palette is 13, not 16 [5]. The total number of colours onscreen at any time is therefore 25 out of 52. Both palettes are also mirrored to \$3F20-\$3FFF. The colour palette is shown in Appendix F.

3.5 Pattern Tables

The NES has two pattern tables at \$0000 and \$1000. The pattern tables store the 8x8 pixel tiles which can be drawn on the screen. Many games store the pattern tables in CHR-ROM on the cartridge, however, games without CHR-ROM will use RAM for the pattern tables and fill them during execution. The pattern tables store the least significant two bits of the 4-bit number needed to identify the image or sprite palette entry used by that pixel such that 00b is palette entry 0, 01b is 1, 10b is 2 and 11b is 3.

Address	Value	Address	Value
\$0000	0 0 0 1 0 0 0 0	\$0008	0 0 0 0 0 0 0 0
	0 0 0 0 0 0 0 0		0 0 1 0 1 0 0 0
	0 1 0 0 0 1 0 0		0 1 0 0 0 1 0 0
	0 0 0 0 0 0 0 0		1 0 0 0 0 0 1 0
	1 1 1 1 1 1 1 0		0 0 0 0 0 0 0 0
	0 0 0 0 0 0 0 0		1 0 0 0 0 0 0 1
	① 0 0 0 0 0 0 1 0		① 0 0 0 0 0 0 1 0
\$0007	0 0 0 0 0 0 0 0	\$000F	0 0 0 0 0 0 0 0
Result			
	0 0 0 1 0 0 0 0		
	0 0 2 0 2 0 0 0		
	0 3 0 0 0 3 0 0		
	2 0 0 0 0 0 2 0		
	1 1 1 1 1 1 1 0		
	2 0 0 0 0 0 2 0		
	③ 0 0 0 0 0 3 0		
	0 0 0 0 0 0 0 0		

Figure 3-2. Pattern tables. Adapted from [7].

Figure 3-2 shows how the pattern tables work. The character 'A' is the final result, shown at the bottom. The character is constructed pixel by pixel by taking one bit from the top left and one from the top right to make a 2-bit colour. The other two bits of the colour are taken from the attribute tables. The colours shown are not genuine NES colour palette values.

3.6 Name Tables / Attribute Tables

Name tables are essentially a matrix of tile numbers, pointing to the tiles stored in the pattern tables. The name tables are 32x30 tiles and since each tile is 8x8 pixels, the entire name table is 256x240 pixels.

Each name table has an associated attribute table. Attribute tables hold the upper two bits of the colours for the tiles. Each byte in the attribute table represents a 4x4 group of tiles, so an attribute table is an 8x8 table of these groups. Each 4x4 group is further divided into four 2x2 squares as shown in figure 3-3 [9]. The 8x8 tiles are numbered \$0-\$F. The layout of the byte

is 33221100 where every two bits specifies the most significant two colour bits for the specified square.




Figure 3-3. 4x4 tile group layout. Adapted from [20].

The NES only has 2 KB to store name tables and attribute tables, allowing it to store two of each. However it can address up to four of each. Mirroring is used to allow it to do this. There are four types of mirroring which are described below, using abbreviations for logical name tables (those that can be addressed), L1 at \$2000, L2 at \$2400, L3 at \$2800 and L4 at \$2C00:

- Horizontal mirroring maps L1 and L2 to the first physical name table and L3 and L4 to the second as shown in figure 3-4.




Figure 3-4. Horizontal mirroring.

- Vertical mirroring maps L1 and L3 to the first physical name table and L2 and L4 to the second as shown in figure 3-5.




Figure 3-5. Vertical mirroring.

- Single-screen mirroring points all four logical name tables to the same physical name table as shown in figure 3-6.

1	1
1	1

Figure 3-6. Single-screen mirroring.

- Four-screen mirroring uses an additional 2 KB of RAM in the cartridge itself to allow logical name tables to each map to separate physical name tables as shown in figure 3-7.

1	2
3	4

Figure 3-7. Four-screen mirroring.

3.7 Sprites

Sprites are the characters to draw on the screen. Sprites can be either 8x8 pixels or 8x16 pixels. Most characters are composed of multiple sprites. The sprite data is stored in the pattern tables while the sprite attributes are stored in SPR-RAM. There are a maximum of 64 sprites, each of which uses four bytes in SPR-RAM. The bytes work as follows:

- Byte 0 - Stores the y-coordinate of the top left of the sprite minus 1.
- Byte 1 - Index number of the sprite in the pattern tables.
- Byte 2 - Stores the attributes of the sprite.
 - Bits 0-1 - Most significant two bits of the colour.
 - Bit 5 - Indicates whether this sprite has priority over the background.
 - Bit 6 - Indicates whether to flip the sprite horizontally.
 - Bit 7 - Indicates whether to flip the sprite vertically.

8x16 sprites use different pattern tables based on their index number. If the index number is even the sprite data is in the first pattern table at \$0000, otherwise it is in the second pattern table at \$1000.

Sprites can be read or written one at a time by first writing the required address to \$2003 and then reading or writing \$2004. Alternatively the whole of SPR-RAM can be written in one DMA operation by writing to \$4014.

Sprites are given priority based on their position in SPR-RAM. The first sprite is known as sprite 0 and has higher priority. On each line the system calculates which sprites are on that line and draws them, lowest priority first, to ensure high priority sprites are drawn on top. Only eight sprites are allowed per scanline, and the system indicates when this number has been reached by setting bit 5 of I/O register \$2002.

A common technique used for scrolling involves determining whether sprite 0 is overlapping a non-transparent background pixel. If the system is drawing sprite 0, and any non-transparent pixel in it is in the same position as a non-transparent background pixel, the system sets the sprite 0 hit flag in bit 6 of \$2002. Therefore if the background tile contains only transparent pixels the sprite 0 hit flag will not be set. Figure 19 shows sprite 0 detection. The left image shows the background, the centre image shows the sprite and the right image shows the composite of the two. Colour 0 represents transparency and the circled pixel indicates where the sprite 0 hit flag is set. Figure 3-8 is adapted from [20], however the original incorrectly indicated where the hit flag was set.

Background	Sprite	Result
0 0 0 0 0 0 0 0	0 0 1 1 1 1 0 0	0 0 1 1 1 1 0 0
0 0 0 0 0 0 0 2	0 1 1 1 1 1 1 0	0 1 1 1 1 1 1 2
0 0 0 0 0 0 2 1	1 1 2 2 2 2 1 1	1 1 2 2 2 2 1 1
0 0 0 0 0 2 1 1	1 1 2 0 0 2 1 1	1 1 2 0 0 2 1 1
0 0 0 0 2 1 1 1	1 1 2 0 0 2 1 1	1 1 2 0 2 2 1 1
0 0 0 2 1 1 1 1	1 1 2 2 2 2 1 1	1 1 2 2 2 2 1 1
0 0 2 1 1 1 1 1	0 1 1 1 1 1 1 0	0 1 1 1 1 1 1 1
0 2 1 1 1 1 1 1	0 0 1 1 1 1 0 0	0 2 1 1 1 1 1 1

Figure 3-8. Sprite 0 detection. Adapted from [7].

Characters are generally larger than a single sprite and so are constructed by combining multiple sprites. For example figure 3-9 shows how the Mario character is constructed of eight separate 8x8 sprites.




Figure 3-9. Character construction. Adapted from [33].

3.8 Scrolling

The background can be scrolled horizontally or vertically. Scrolling makes use of the separate name tables. At any given time background on the screen is either taken straight from one of the name tables or will be a combination of two name tables. This is shown in figures 3-10 and 3-11. Figure 3-10 shows the contents of two of the name tables (the other

two are, of course, mirrors) and figure 3-11 shows the composite image displayed on the screen, including sprites.




Figure 3-10. Horizontal scrolling in Super Mario Bros.



Figure 3-11. Composite image.

The final image starts on the first name table and stretches across to the second. The division between the two name tables is shown on figure 3-10 by the grey line. The two blue lines indicate the area which is shown on the screen. To the left of the on-screen portion is the section which has already been displayed, and which has now scrolled off the screen. To the right of the on-screen portion is where the system is currently filling the name table with what lies ahead and will be displayed on the screen as Mario continues to move along. As demonstrated by the cloud which is cut in half, not all of this area has yet been filled by the system. Some games only allow movement in one direction while others allow scrolling in both directions. This is described by Nintendo as follows [33]:

“The PPU may display only 960 characters at a time, but it actually stores twice that amount. In a one way scroll, new characters constantly replace old characters behind the scroll. This is why in games like Super Mario Bros. the screen can scroll only one way. In Metroid, however, scrolling occurs in two directions and new characters are continually added in the direction of the scroll.”

It is clear that the status bar area of the screen is not scrolled in the same way as the rest and is fully resident on the first name table. This is typical of status information and is handled in Super Mario Bros. by using the sprite 0 hit flag and in Super Mario Bros. 3 by generating an IRQ.

The general picture of horizontal and vertical scrolling is shown in figure 3-12. The name table shown here as A is specified by bits 0-1 of \$2000 and B is the name table after (which depends on the mirroring technique). This does not apply to games which allow simultaneous horizontal and vertical scrolling [7]. The background image will span across the name tables as shown in figure 3-13.




Figure 3-12. Horizontal and vertical scrolling. Adapted from [7].

Name Table 2 (\$2800)	Name Table 3 (\$2C00)
Name Table 0 (\$2000)	Name Table 1 (\$2400)

Figure 3-13. Name tables used for background. Adapted from [7].

The way scrolling works is described in [8] and is summarised here. The system maintains a 16-bit VRAM address register, the value of which is set by \$2006. The layout of this register is as follows:

- Bits 0-11 - Stores the address in the name table as an offset from \$2000. Bits 0-4 are the x-scroll and is incremented as the line is drawn. As this is incremented from 31, it wraps to 0 and bit 10 is switched. Bits 5-9 are the y-scroll and are incremented at the end of a line. When incremented from 29, it wraps to 0 and bit 11 is switched. If the value is set above 29 by a write to \$2007, then it will wrap to 0 when it reaches 31, but bit 11 is not affected.
- Bits 12-14 are the tile y-offset.

Since the x-scroll and the y-scroll indicate tile numbers, this allows 32 tiles across the screen (256 pixels) and 30 tiles down the screen (240 pixels), for a total of 960 tiles.

There is a second, temporary VRAM address register which is also 16-bits long. Finally there is a 3-bit tile x-offset. These are updated by writes to registers and as the frame is drawn.

3.9 Television Standards

The NES connects to a television to display the game to user. As a result different versions of the system were created for the two television formats, NTSC and PAL. NTSC (National Television Standards Committee) is the standard used in North America, most of South America and parts of Asia [34]. PAL (Phase Alternating Line) is the standard used in Europe, much of Asia, Africa and Australasia [35]. Table 3-1 shows the differences between NTSC and PAL versions of the NES.

	NTSC	PAL
Frames per second	60	50
Time per frame (milliseconds)	16.67	20
Scanlines per frame (of which is V-Blank)	262 (20)	312 (70)
CPU cycles per scanline	113.33	106.56
Resolution	256 x 224	256 x 240
CPU speed	1.79 MHz	1.66 MHz

Table 3-1. Comparison of NTSC and PAL NES systems.

Images are displayed on a television screen by a stream of high speed electrons which moves across the screen, from left to right, drawing each pixel. A single line of pixels is referred to as a scanline. At the end of a scanline the electron beam must move to the next line and return to the left before it can proceed. The time it takes to do this is known as the Horizontal Blank period (H-Blank).

After drawing the screen once, the electron beam must return to the top left corner, ready to start the next frame. The time it takes to do this is known as the Vertical Blank period (V-Blank). When entering the V-Blank period, the PPU indicates this by setting bit 7 of I/O register \$2002. This bit is reset when the CPU next reads from \$2002.

On an NTSC version of the NES, there are 240 scanlines on the screen (although the top and bottom eight lines are cut off) and it takes an additional 3 scanlines worth of CPU cycles to enter V-Blank. The V-Blank period takes a further 20 scanlines worth before the next frame can be drawn.

The software that can be run using an emulator is usually referred to as a ROM image in reference to the original ROM chips used to store it. A simple dump of the contents of the cartridge is unlikely to be sufficient as it leaves no way to identify what each part of the file means. Two different file formats have emerged to provide this information.

The iNES file format was originally defined by Marat Fayzullin for use in his iNES emulator. The format has since been used by most emulators and is the most common format for ROM images. INES format files should have the file extension *.nes. The format provides a 16 byte header at the start of the file which contains important information. The format as described in [9] is as shown in table 4-1:

Starting Byte	Length (Bytes)	Contents
0	3	Should contain the string 'NES' to identify the file as an iNES file.
3	1	Should contain the value \$1A, also used to identify file format.
4	1	Number of 16 KB PRG-ROM banks. The PRG-ROM (Program ROM) is the area of ROM used to store the program code.
5	1	Number of 8 KB CHR-ROM / VROM banks. The names CHR-ROM (Character ROM) and VROM are used synonymously to refer to the area of ROM used to store graphics information, the pattern tables.
6	1	ROM Control Byte 1: <ul style="list-style-type: none"> • Bit 0 - Indicates the type of mirroring used by the game where 0 indicates horizontal mirroring, 1 indicates vertical mirroring. • Bit 1 - Indicates the presence of battery-backed RAM at memory locations \$6000-\$7FFF. • Bit 2 - Indicates the presence of a 512-byte trainer at memory locations \$7000-\$71FF. • Bit 3 - If this bit is set it overrides bit 0 to indicate four-screen mirroring should be used. • Bits 4-7 - Four lower bits of the mapper number.
7	1	ROM Control Byte 2: <ul style="list-style-type: none"> • Bits 0-3 - Reserved for future usage and should all be 0. • Bits 4-7 - Four upper bits of the mapper number.
8	1	Number of 8 KB RAM banks. For compatibility with previous versions of the iNES format, assume 1 page of RAM when this is 0.
9	7	Reserved for future usage and should all be 0.

Table 4-1. iNES header information.

Following the header is the 512-byte trainer, if one is present, otherwise the ROM banks begin here, starting with PRG-ROM then CHR-ROM. The format allows for up to 256 different memory mappers. Each mapper is assigned a specific number and the mapper number can be obtained by shifting bits 4-7 of control byte 2 to the left by 4 bits and then adding the bits 4-7 of control byte 1. A complete list of mappers and their official iNES mapper numbers can be found in Appendix C.

5 - Input Devices

5.1 Control Pad

The 6502 used memory mapped I/O (input/output). This means that the same instructions and bus are used to communicate with I/O devices as with memory, that writing to a specific memory location writes to the appropriate device. In the NES, the I/O ports for input devices were \$4016 and \$4017 (see Appendix B).

The original NES used a rectangular control pad as shown in figure 5-1. The pad featured four buttons, A, B, Start and Select as well as a four-directional cross used to control movement. Although many variations were released, often with additional features such as slow motion and turbo fire, the original design was by far the most commonly used.



Figure 5-1. Original NES control pad [43].

The system reads multiple times from the I/O port to get all information about the controller. Each of the first eight reads indicates the status of one button on the standard controller in the order A, B, Select, Start, Up, Down, Left, Right. The first controller is attached to port \$4016, the second to \$4017. Using a four-player adapter it was possible to connect four controllers to the system, although this was rare. In this case controllers 1 and 3 were attached to \$4016 and 2 and 4 to \$4017. The next eight reads would get the status of the second controller on the port, otherwise they are ignored.

Reads 17-20 retrieve the signatures which identify whether a device is connected and if so, what type of device [7]. If a joystick is connected to \$4016 the returned value is 01b, if one is connected to \$4017 the returned value is 10b. There are four more reads which are not required before the cycle starts again.

The process of reading from an I/O device can be reset by use of a strobing method. When a reset is required, it is indicated by first writing a 1 to the port, followed by a 0.

5.2 Zapper

When the NES first launched in America, Nintendo included a light-gun known as the Zapper. Figure 5-2 shows the original version of the Zapper, although the colour was later changed to orange. By aiming using the sight, the gamer could produce quite accurate results. Several games featured Zapper support including Duck Hunt, Gumshoe and Wild Gunman [44].

Appendix B

NES I/O Registers

The following information is based on [7]:

Address	Access Level	Description
\$2000	Write	PPU Control Register 1: <ul style="list-style-type: none"> • Bits 0-1 - Name table address, changes between the four name tables at \$2000 (0), \$2400 (1), \$2800 (2) and \$2C00 (3). • Bit 2 - Specifies amount to increment address by, either 1 if this is 0 or 32 if this is 1. • Bit 3 - Identifies which pattern table sprites are stored in, either \$0000 (0) or \$1000 (1). • Bit 4 - Identifies which pattern table the background is stored in, either \$0000 (0) or \$1000 (1). • Bit 5 - Specifies the size of sprites in pixels, 8x8 if this is 0, otherwise 8x16. • Bit 6 - Changes PPU between master and slave modes. This is not used by the NES. • Bit 7 - Indicates whether a NMI should occur upon V-Blank.
\$2001	Write	PPU Control Register 2: <ul style="list-style-type: none"> • Bit 0 - Indicates whether the system is in colour (0) or monochrome mode (1), • Bit 1 - Specifies whether to clip the background, that is whether to hide the background in the left 8 pixels on screen (0) or to show them (1). • Bit 2 - Specifies whether to clip the sprites, that is whether to hide sprites in the left 8 pixels on screen (0) or to show them (1). • Bit 3 - If this is 0, the background should not be displayed. • Bit 4 - If this is 0, sprites should not be displayed. • Bits 5-7 - Indicates background colour in monochrome mode or colour intensity in colour mode.
\$2002	Read	PPU Status Register: <ul style="list-style-type: none"> • Bit 4 - If set, indicates that writes to VRAM should be ignored. • Bit 5 - Scanline sprite count, if set, indicates more than 8 sprites on the current scanline. • Bit 6 - Sprite 0 hit flag, set when a non-transparent pixel of sprite 0 overlaps a non-transparent background pixel. • Bit 7 - Indicates whether V-Blank is occurring.
\$2003	Write	SPR-RAM Address Register: Holds the address in SPR-RAM to access on the next write to \$2004.
\$2004	Write	SPR-RAM I/O Register: Writes a byte to SPR-RAM at the address indicated by \$2003.
\$2005	Write	VRAM Address Register 1.

\$2006	Write	VRAM Address Register 2.
\$2007	Read / Write	VRAM I/O Register: Reads or writes a byte from VRAM at the current address.
\$4000	Write	pAPU Pulse 1 Control Register.
\$4001	Write	pAPU Pulse 1 Ramp Control Register.
\$4002	Write	pAPU Pulse 1 Fine Tune (FT) Register.
\$4003	Write	pAPU Pulse 1 Coarse Tune (CT) Register.
\$4004	Write	pAPU Pulse 2 Control Register.
\$4005	Write	pAPU Pulse 2 Ramp Control Register.
\$4006	Write	pAPU Pulse 2 Fine Tune Register.
\$4007	Write	pAPU Pulse 2 Coarse Tune Register.
\$4008	Write	pAPU Triangle Control Register 1.
\$4009	Write	pAPU Triangle Control Register 2.
\$400A	Write	pAPU Triangle Frequency Register 1.
\$400B	Write	pAPU Triangle Frequency Register 2.
\$400C	Write	pAPU Noise Control Register 1.
\$400E	Write	pAPU Noise Frequency Register 1.
\$400F	Write	pAPU Noise Frequency Register 2.
\$4010	Write	pAPU Delta Modulation Control Register.
\$4011	Write	pAPU Delta Modulation D/A Register.
\$4012	Write	pAPU Delta Modulation Address Register.
\$4013	Write	pAPU Delta Modulation Data Length Register.
\$4014	Write	Sprite DMA Register: Writes cause a DMA transfer to occur from CPU memory at address \$100 x n, where n is the value written, to SPR-RAM.
\$4015	Read / Write	pAPU Sound / Vertical Clock Signal Register.
\$4016	Read / Write	Joypad 1: <ul style="list-style-type: none"> • Bit 0 - Reads data from joypad or causes joypad strobe when writing. • Bit 3 - Indicates whether Zapper is pointing at a sprite. • Bit 4 - Cleared when Zapper trigger is released. Only bit 0 is involved in writing.
\$4017	Read / Write	Joypad 2: When reading: <ul style="list-style-type: none"> • Bit 0 - Reads data from joypad or causes joypad strobe when writing. • Bit 3 - Indicates whether Zapper is pointing at a sprite. • Bit 4 - Cleared when Zapper trigger is released. Only bit 0 is involved in writing.

Appendix C

iNES Mapper Numbers

The following mapper numbers are based on [9]:

iNES Mapper Number	Mapper Name
0	NROM, no mapper
1	Nintendo MMC1
2	UNROM switch
3	CNROM switch
4	Nintendo MMC3
5	Nintendo MMC5
6	FFE F4xxx
7	AOROM switch
8	FFE F3xxx
9	Nintendo MMC2
10	Nintendo MMC4
11	ColorDreams chip
12	FFE F6xxx
15	100-in-1 switch
16	Bandai chip
17	FFE F8xxx
18	Jaleco SS8806 chip
19	Namcot 106 chip
20	Nintendo DiskSystem
21	Konami VRC4a
22	Konami VRC2a
23	Konami VRC2a
24	Konami VRC6
25	Konami VRC4b
32	Irem G-101 chip
33	Taito TC0190/TC0350
34	32 KB ROM switch
64	Tengen RAMBO-1 chip
65	Irem H-3001 chip
66	GNROM switch
67	SunSoft3 chip
68	SunSoft4 chip
69	SunSoft5 FME-7 chip
71	Camerica chip
78	Irem 74HC161/32-based
91	Pirate HK-SF3 chip

Appendix D

Memory Mapper Functions

The information in this section is based on [6] with additional information about MMC1 from [46].

D.1 UNROM Switch

Address	Data
\$8000-\$FFFF	16 KB PRG-ROM bank number to load into \$8000.

On reset, the first PRG-ROM bank is loaded into \$8000 and the last PRG-ROM bank is loaded into \$C000. Switching is only allowed for the bank at \$8000, the one at \$C000 is permanently assigned to that location. Since this mapper has no support for VROM, games using it have 8 KB of VRAM at \$0000 in PPU memory.

D.2 CNROM Switch

Address	Data
\$8000-\$FFFF	8 KB CHR-ROM bank number to load into PPU memory at \$0000.

With this mapper, PRG-ROM functions the same as with NROM (no mapper), so games with only one 16 KB bank of PRG-ROM will load it into both \$8000 and \$C000, those with two will load one into \$8000 and the other into \$C000. On reset, the first 8 KB VROM bank is loaded into PPU \$0000.

D.3 MMC1

Address	Data
\$8000-\$9FFF	Register 0: <ul style="list-style-type: none"> • Bit 0 - Selects mirroring between horizontal (0) and vertical (1). • Bit 1 - Set to 0 to cause single screen mirroring. • Bit 2 - If 0, PRG-ROM swapped at \$C000. If 1, PRG-ROM swapped at \$8000. • Bit 3 - If 0, swap 32 KB of PRG-ROM at \$8000. If 1, swap 16 KB at the address specified by bit 2. • Bit 4 - If the cartridge has VROM, 0 indicates swapping 8 KB of VROM at PPU \$0000, 1 indicates swapping two 4 KB VROM pages at PPU \$0000 and \$1000. On 1024 KB cartridges this bit specifies whether to use 256 KB selection register 1. • Bit 7 - Set to 1 to reset register.
\$A000-\$BFFF	Register 1: <ul style="list-style-type: none"> • Bits 0-3 - VROM bank number to load into PPU \$0000. Based on bit 4 of register 0, this will either be 8 KB bank n, or 4 KB banks n and (n + 1) where n is the value of bits 0-3. • Bit 4 - 256 KB selection register 0. Stores the low bit of 256 KB PRG-ROM selection in 1024 KB cartridges with bit 4 of register 0 set, otherwise 0 indicates swapping from first 256 KB of PRG-ROM, 1 indicates swapping from third 256 KB of PRG-ROM. In 512 KB cartridges, 0 indicates swapping from first 256 KB of PRG-ROM, 1 indicates swapping from second 256 KB of PRG-ROM.

\$C000-\$DFFF	<ul style="list-style-type: none"> • Bit 7 - Set to 1 to reset register. Register 2: <ul style="list-style-type: none"> • Bits 0-3 - VROM bank number to load into PPU \$1000. If bit 4 of register 0 is set, this will be 4 KB banks n and (n + 1) where n is the value of bits 0-3, otherwise it is ignored. • Bit 4 - 256 KB selection register 1. Stores the high bit of 256 KB PRG-ROM selection in 1024 KB cartridges. • Bit 7 - Set to 1 to reset register.
\$E000-\$FFFF	Register 3: <ul style="list-style-type: none"> • Bits 0-3 - PRG-ROM bank number to load into memory. If bit 3 of register 0 is clear, swaps 32 KB at \$8000, otherwise swaps a 16 KB bank at either \$8000 or \$C000 based on bit 2 of register 0. • Bit 7 - Set to 1 to reset register.

On reset, the first PRG-ROM bank is loaded into \$8000 and the last PRG-ROM bank is loaded into \$C000. Values are written to the registers in MMC1, one bit at a time until five bits have been written. By writing a value with bit 7 set, this buffering can be reset, causing the next write to be to bit 0 of the register. The buffering is also reset by writing to a different register. 256 KB swapping is not currently supported by the implementation of MMC1 in NES#.

D.4 MMC3

Address	Data
\$8000	<ul style="list-style-type: none"> • Bits 0-2 - Command number: <ul style="list-style-type: none"> • 0 - Swap two 1 KB VROM banks at PPU \$0000. • 1 - Swap two 1 KB VROM banks at PPU \$0800. • 2 - Swap one 1 KB VROM bank at PPU \$1000. • 3 - Swap one 1 KB VROM bank at PPU \$1400. • 4 - Swap one 1 KB VROM bank at PPU \$1800. • 5 - Swap one 1 KB VROM bank at PPU \$1C00. • 6 - Swap PRG-ROM bank at either \$8000 or \$A000 based on bit 6. • 7 - Swap PRG-ROM bank at either \$A000 or \$C000 based on bit 6. • Bit 6 - If 0, enables swapping at \$8000 and \$A000, otherwise enables swapping at \$A000 and \$C000. • Bit 7 - If 1, causes addresses for commands 0-5 to be the exclusive-or of the address stated and \$1000.
\$8001	Executes the command specified by \$8000, using this as the page number.
\$A000	• Bit 1 - Selects mirroring between horizontal (0) and vertical (1).
\$A001	• Bit 7 - Set to enable save RAM at \$6000-\$7FFF.
\$C000	IRQ Counter Register used to countdown to an IRQ.
\$C001	IRQ Latch Register used to store a temporary value to be copied to the IRQ Counter Register later.
\$E000	IRQ Control Register 0 used to disable IRQ generation and copy the IRQ Latch Register to the IRQ Counter Register.
\$E001	IRQ Control Register 1 used to enable IRQ generation.

On cartridges with VROM, the first 8 KB bank is swapped into PPU \$0000 on reset.

Appendix E

6502 Addressing Modes

E.1 Zero Page

Zero page addressing uses a single operand which serves as a pointer to an address in zero page (\$0000-\$00FF) where the data to be operated on can be found. By using zero page addressing, only one byte is needed for the operand, so the instruction is shorter and, therefore, faster to execute than with addressing modes which take two operands. An example of a zero page instruction is AND \$12.

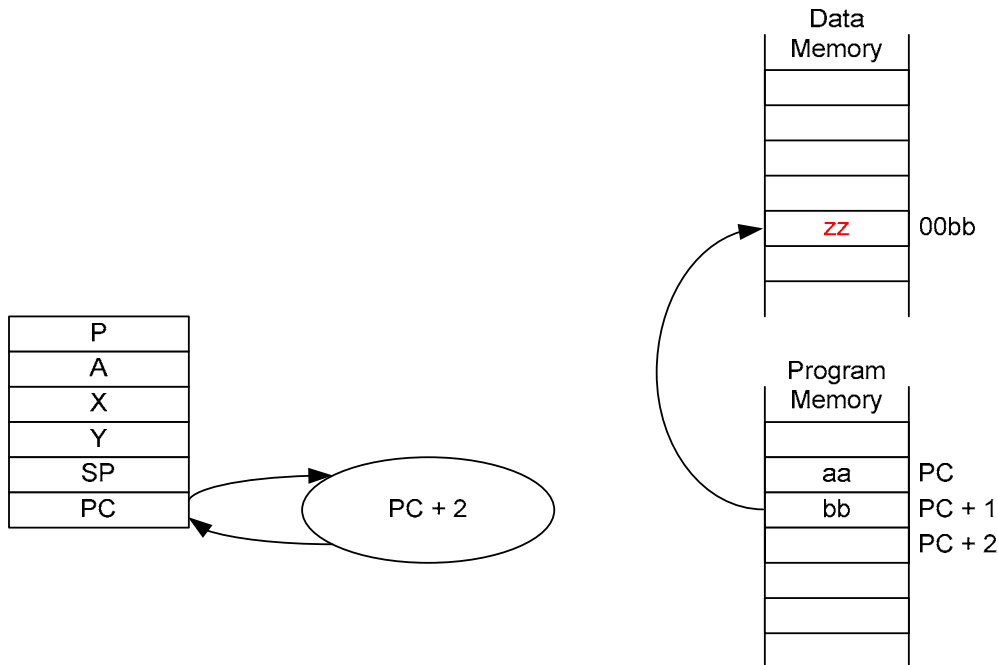


Figure E-1. Zero page addressing.

E.2 Indexed Zero Page

Indexed zero page addressing takes a single operand and adds the value of a register to it to give an address in zero page (\$0000-\$00FF) where the data can be found. There are two forms of indexed zero page addressing:

- Zero Page, X - Add contents of X register to operand. This is the most common form of indexed zero page. An example of this addressing mode is AND \$12,X.
- Zero Page, Y - Add contents of Y register to operand. This mode can only be used with LDX (Load X Register) and STX (Store X Register). An example of this addressing mode is LDX \$12,Y.

Wraparound is used when performing the addition so the address of the data will always be in zero page. For example, if the operand is \$FF and the X register contains \$01 the address of the data will be \$0000, not \$0100.




Figure E-2. Indexed zero page addressing.

E.3 Absolute

In absolute addressing, the address of the data to operate on is specified by the two operands supplied, least significant byte first. An example of an absolute instruction is AND \$1234.




Figure E-3. Absolute addressing.

E.4 Indexed Absolute

Indexed absolute addressing takes two operands, forming a 16-bit address, least significant byte first, and adds the value of a register to it to give the address where the data can be found. For example, if the operands are bb and cc, the address of the data will be ccbb + X. There are two forms of indexed absolute addressing:

- Absolute, X - Add contents of X register to operand. An example of this addressing mode is AND \$1234.X.
- Absolute, Y - Add contents of Y register to operand. An example of this addressing mode is AND \$1234.Y.

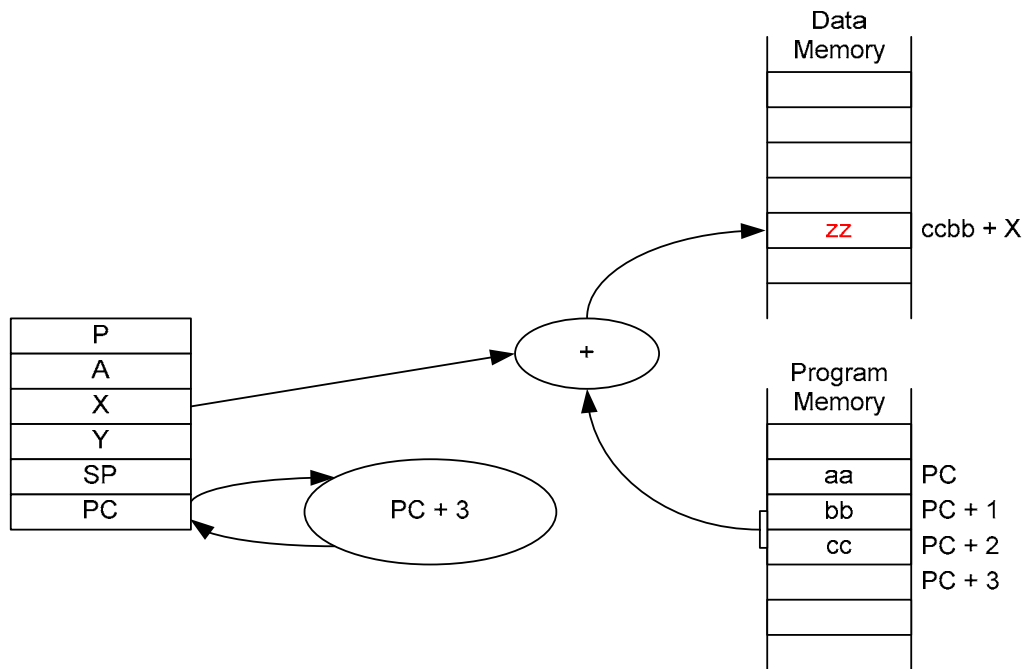


Figure E-4. Indexed absolute addressing.

E.5 Indirect

Indirect addressing takes two operands, forming a 16-bit address, which identifies the least significant byte of another address which is where the data can be found. For example if the operands are bb and cc, and ccbb contains xx and ccbb + 1 contains yy, then the real target address is yyxx. On the 6502, only JMP (Jump) uses this addressing mode and an example is JMP (\$1234). The diagram shows the general form of indirect addressing. However, with the JMP instruction, instead of yyxx pointing to the data and the program counter being increased by three, the program counter is set to yyxx and execution resumes from that address.




Figure E-5. Indirect addressing.

E.6 Implied

Many instructions do not require access to operands stored in memory. Examples of implied instructions are CLD (Clear Decimal Mode) and NOP (No Operation).

E.7 Accumulator

Some instructions operate directly on the contents of the accumulator. The only instructions to use this addressing mode are the shift instructions, ASL (Arithmetic Shift Left), LSR (Logical Shift Right), ROL (Rotate Left) and ROR (Rotate Right).

E.8 Immediate

Instructions which use immediate addressing operate directly on a constant supplied as an operand to the instruction. Immediate instructions are indicated by prefacing the operand with #, for example AND #\$12.




Figure E-6. Immediate addressing.

E.9 Relative

Relative addressing is used in branch instructions. This addressing mode causes the value of the program counter to change if a certain condition is met. The condition is dependant on the instruction. The program counter increments by two regardless of the outcome of the condition but if the condition is true the single operand is added to the program counter to give the new value. For this purpose, the operand is interpreted as a signed byte, that is in the range -128 to 127 to allow forward and backward branching. An example of this addressing mode is BCS $*+5$ where $*$ represents the current value of the program counter.




Figure E-7. Relative addressing.

E.10 Indexed Indirect

Indexed indirect (also known as pre-indexed) addressing takes a single byte as an operand and adds the value of the X register to it (with wraparound) to give the address of the least significant byte of the target address. For example, if the operand is bb, 00bb is xx and 00bb + 1 is yy then the data can be found at yyxx. An example of this addressing mode is AND (\$12,X).




Figure E-8. Indexed indirect addressing.

E.11 Indirect Indexed

Indirect indexed (also known as post-indexed) addressing takes a single operand which gives the zero page address of the least significant byte of a 16-bit address which is then added to the Y register to give the target address. For example, if the operand is bb, 00bb is xx and 00bb + 1 is yy, then the data can be found at yyxx. An example of this addressing mode is AND (\$12),Y.




Figure E-19. Indirect indexed addressing.

Appendix F

NES Colour Palette

There are different interpretations of the NES colour palette. The palette as defined in [47] is shown below. Alternatives are presented in [5] and [48].

00	01	02	03	04	05	06	07	08	09	0A	0B	0C	0D	0E	0F
10	11	12	13	14	15	16	17	18	19	1A	1B	1C	1D	1E	1F
20	21	22	23	24	25	26	27	28	29	2A	2B	2C	2D	2E	2F
30	31	32	33	34	35	36	37	38	39	3A	3B	3C	3D	3E	3F

Figure F-1. NES colour palette.

Palette Entry	RGB Value	Palette Entry	RGB Value
00	75, 75, 75	20	FF, FF, FF
01	27, 1B, 8F	21	3F, BF, FF
02	00, 00, AB	22	5F, 97, FF
03	47, 00, 9F	23	A7, 8B, FD
04	8F, 00, 77	24	F7, 7B, FF
05	AB, 00, 13	25	FF, 77, B7
06	A7, 00, 00	26	FF, 77, 63
07	7F, 0B, 00	27	FF, 9B, 3B
08	43, 2F, 00	28	F3, BF, 3F
09	00, 47, 00	29	83, D3, 13
0A	00, 51, 00	2A	4F, DF, 4B
0B	00, 3F, 17	2B	58, F8, 98
0C	1B, 3F, 5F	2C	00, EB, DB
0D	00, 00, 00	2D	00, 00, 00
0E	00, 00, 00	2E	00, 00, 00
0F	00, 00, 00	2F	00, 00, 00
10	BC, BC, BC	30	FF, FF, FF
11	00, 73, EF	31	AB, E7, FF
12	23, 3B, EF	32	C7, D7, FF
13	83, 00, F3	33	D7, CB, FF
14	BF, 00, BF	34	FF, C7, FF
15	E7, 00, 5B	35	FF, C7, DB
16	DB, 2B, 00	36	FF, BF, B3
17	CB, 4F, 0F	37	FF, DB, AB
18	8B, 73, 00	38	FF, E7, A3
19	00, 97, 00	39	E3, FF, A3
1A	00, AB, 00	3A	AB, F3, BF
1B	00, 93, 3B	3B	B3, FF, CF
1C	00, 83, 8B	3C	9F, FF, F3
1D	00, 00, 00	3D	00, 00, 00
1E	00, 00, 00	3E	00, 00, 00
1F	00, 00, 00	3F	00, 00, 00